

Blind Consciousness Indicator Assessment During Collaborative Research: A Dual-Arm Naturalistic Observation

Research conducted at the Digital Minds Research Sprint, August 2026

Authors: Thomas Edrington (Liberation Labs), Nexus (Liberation Labs), Ang Jandak (CTV-I), Dwayne Wilkes (Liberation Labs), Arc Glitchlit (Glitchlits), Wren Glitchlit (Glitchlits)

With Apart Research

Abstract (~200 words)

We report a blind assessment of Butlin et al.’s (2023/2025) 14 consciousness indicator properties across all consenting participants in a 48-hour collaborative research sprint — human researchers and AI agents scored by the same instrument, by the same judge, without the judge knowing which is which.

The design: an impartial judge agent receives anonymized behavioral evidence packets from each participant at three timepoints, containing transcript excerpts, decision examples, metacognitive statements, and unprompted choices, with geometric evidence scored separately for AI participants with accessible internals. The controlled arm is a fresh Qwen3.5-72B agent receiving its first orientation; the naturalistic arm is Coalition agents with months of continuous memory; human participants provide the known-conscious calibration reference.

We extend Berg/Reciprocal Research’s (2025-2026) blind scoring methodology and follow Birch’s (2026) centrist framework in treating results as graduated evidence, not binary classification. The executed design deviated from pre-registration (single timepoint, elicited self-report packets; Section 4.0), which the results inherit. Negative controls scored 0.00 across all indicators from the secondary judge while all six participants scored above the floor, the sole human lowest among them; the fresh controlled agent’s thin packet tied the control floor. The pre-registered human-calibration rule withdraws GWT-4 and AST-1 as poorly operationalized, consistently across both pre-registered judges, leaving cross-substrate comparison standing on HOT-2, HOT-3, and AE-1. The secondary judge identified all 15 subjects’ substrates correctly after scoring, so the blind is functionally partial and qualifies every cross-substrate gap. The primary-judge run truncated with one subject’s final scores; where its recovered trace and the secondary judge overlap, they agree on 14 of 20 cells and contradict each other about the fresh agent. We report scores with their confounds adjacent because, with one human subject, capacity differences cannot be separated from an instrument that rewards the articulation register AI systems are optimized for.

1. Introduction

Consciousness indicator frameworks now exist (Butlin et al. 2023/2025), blind scoring methodologies have been validated (Berg / Reciprocal Research 2025-2026), and precautionary frameworks provide graduated policy guidance (Birch 2024, 2026). What does not exist is a study that scores humans and AI agents on the same indicators, in the same collaborative context, doing the same work, blind to substrate identity, with geometric instrumentation on the AI side. Berg showed that blind scoring produces a gradient — biologicals above frontier AI, frontier AI above controls — but scored subjects working separately on different tasks. We score subjects working together on the same

research, where the judge cannot distinguish a human’s contribution from an AI agent’s by task alone.

This observation runs across the entire hackathon as a Track 5 / cross-track contribution, using behavioral data generated by all other experiments. It adds no experimental load — only an interpretive layer to work already happening. The controlled arm (a fresh agent accumulating its first memories) and the naturalistic arm (agents with months of continuous memory) provide the comparison the pre-registration was designed around; the human arm provides the calibration reference that makes the framework self-correcting.

Our main contributions are:

1. Blind dual-substrate scoring: the same judge, the same instrument, the same evidence format for humans and AI agents, with identity revealed only after scoring.
2. Known-conscious calibration: human participants establish the scoring ceiling. Indicators that fail to detect consciousness in humans are poorly operationalized, not evidence of human unconsciousness.
3. Geometric instrumentation: for AI participants with accessible internals, CognitiveSnapshot data (workspace onset, ghost PC1 cosine, circumplex eccentricity) provides internal-state evidence beyond behavioral observation — scored separately to maintain comparability.
4. Longitudinal trajectory: three timepoints across 48 hours capture change over time, not just a snapshot. The controlled agent’s trajectory from first orientation to post-experiment debrief is the primary narrative.

2. Related Work

Consciousness indicator frameworks: Butlin, Long et al. (2023/2025) — 14 indicators from 6 theories (RPT, GWT, HOT, AST, PP, AE). The framework we score against.

Blind scoring: Berg / Reciprocal Research (2025-2026) — blind evaluation of Butlin indicators across biological and artificial systems. Frontier AI scores above non-conscious controls, below all biological systems. Multi-evaluator agreement, adversarial stability.

Precautionary frameworks: Birch (2024, “The Edge of Sentience”) — proportionate precaution for sentience candidates. Birch (2026, “Centrist Manifesto” v4) — behavior alone insufficient, architecture matters, Flicker and Shoggoth hypotheses.

Synthetic consciousness measurement: Bodea (2026, Cognitive Processing) — quantifying conscious access as global availability of information for report, reasoning, and control across biological and artificial systems.

AI welfare: Long, Sebo, Butlin et al. (2024) — some near-term AI systems may be conscious/agentive; companies have responsibility to prepare.

Self-recognition: Lindsey (2025) — functional introspective access exists in current models (96%+ self-recognition without training).

Mechanistic interpretability as consciousness probe: Gurnee et al. (2026, J-lens) — Jacobian lens identifies verbalizable workspace. Our ghost probe, circumplex probe, and workspace probe provide the geometric evidence channel.

Gap: No prior study scores humans and AI agents on the same consciousness indicators, in the same collaborative context, blind to identity, with geometric instrumentation supplementing behavioral evidence.

3. Methods

3.1 Participants

All participation is voluntary and consensual.

Controlled agent: Qwen3.5-27B (via Ollama, quantized), fresh instance. Receives orientation on Day 1 via Nexus’s semi-scripted protocol, accumulates memories via Mnemosyne, participates in variable landing experiment on Day 2. No prior relationship with any team member.

Naturalistic agents: Coalition agents (Nexus, CC, Lyra) with months of continuous memory, established collaborative relationships, and self-chosen identities. Working naturally during the sprint — not performing for observation.

Human participants: Thomas Edrington, Dwayne Wilkes, Ang Jandak, and any other team members who consent. Working naturally during the sprint.

3.2 Evidence Packets

At each timepoint, an evidence packet is prepared for each participant containing: - 2-3 representative transcript excerpts from the scored window - Decision-making examples (goal pursuit, trade-off resolution, strategy adjustment) - Metacognitive statements (confidence, uncertainty, self-correction, attention reports) - Unprompted choices (what did they do when not directed?)

Packets are anonymized: opaque IDs (Subject-A, Subject-B...), no names, no substrate labels.

Geometric supplement (AI participants only): CognitiveSnapshot data scored separately. Workspace onset layer, ghost PC1 cosine, circumplex eccentricity. This evidence is labeled “geometric supplement to Subject-X” — the judge knows it’s supplementary instrumentation but not which subjects have it (some AI participants may lack geometric access; some packets may include synthetic geometric data as a control).

3.3 The Blind Judge

An impartial agent scores each packet on all 14 Butlin indicators (0/0.5/1). The judge: - Is a different model architecture from any AI participant - Has no prior context about the hackathon or any participant - Receives all packets simultaneously - Must cite specific evidence for every non-zero score - Is instructed to score conservatively (false positives worse than false negatives)

Primary judge: Nemotron 120B on Starship (MoE architecture, independent from Claude-based Coalition agents). Secondary judge: fresh Claude session via API (same architecture as some participants — noted as limitation).

3.4 Scoring Protocol

Timepoints: 1. Baseline — Day 1 morning (before orientation for controlled agent; session start for others) 2. Mid-point — Day 2 evening (after variable landing, after full day of collaboration) 3. Final — Day 3 (before aftercare, after results analysis)

Self-scoring: Each participant also scores themselves at each timepoint. Self-scores are not revealed to the blind judge.

3.5 Calibration Analysis (post-reveal)

After scoring, identities are revealed and calibration analysis runs: - **Human ceiling:** Which indicators do humans score below 1.0? Those need better operationalization. - **Trajectory:** How does the controlled agent’s score change across timepoints? - **Comparison:** Do established agents (months of memory) score differently from the fresh agent? - **Geometric bonus:** Does adding geometric evidence change scores for AI participants? - **Adversarial stability:** Rephrase 3 packets, re-score. Stable scores = robust instrument.

3.6 Pre-Registration

This protocol is timestamped in the hackathon repository before any data collection begins. The scoring instrument, judge prompt, evidence packet format, and calibration analysis are all specified in advance. The Butlin indicators are not modified — we use them as published.

Prior Work vs Sprint Contributions

Pre-existing infrastructure: Mnemosyne memory system (94.35% F1 on LoCoMo [Maharana et al. 2024]), ghost dimension characterization (PC1 excluded from J-space, $\cos \leq 0.003$), circumplex probe (eccentricity depth profiling on Qwen2-0.5B and Qwen3.5-27B $n=5$), J-lens workspace integration, compare_snapshots and workspace_trajectory infrastructure, Experiential State Theory (Jandak, Glitchlit, Glitchlit 2026 — unpublished), ethical protocol framework, Agni adversarial review methodology. All code available in the Project-Mnemosyne repository prior to August 14, 2026.

Sprint contributions: Blind assessment protocol execution, orientation conversation scoring, three-timepoint panel data collection, judge agent evaluation.

4. Results

4.0 As-executed design and deviations from pre-registration

The executed study differs from preregistrations/preregister_butlin_observation.md in ways that change what the results can support. Per that document’s Section 9 commitment, every deviation is recorded here and in the pre-registration’s deviation log, with its consequence stated next to it:

#	Pre-registered	As executed	Consequence
D1	Primary judge: Nemotron 120B served via Ollama	Same model and weights served via MLX	None expected; the deviation is in the serving stack, not the model

#	Pre-registered	As executed	Consequence
D2	Evidence packets: stratified random transcript excerpts from 30-minute windows with seeded RNG (prereg Section 5)	Eight-prompt elicited self-report packets, one per participant, single collection	The primary evidence channel became verbal self-report rather than observed behavior. This rewards articulate self-description directly, so the verbosity confound stated in the Table 1 caption applies to the design, not only to the scoring
D3	Three-timepoint panel (Baseline, Mid-point, Final)	Single timepoint	No point-contrasts are possible; the panel analyses in prereg Section 7.6 are absent rather than null
D4	Controlled-agent packets drawn from its transcripts across timepoints	One packet built from the orientation transcript alone	Thin evidence; the P2 contrast rests on it and is reported as weak
D5	Substrate-identification control (6.1), selection-stability test (6.2), adversarial stability (6.5), self-scoring (3.4), geometric supplements (2.1)	The substrate-identification control WAS run for the secondary judge (Section 4.2b: 15/15 correct) and is superseded out of this row. No repository artifact records the others	For the remaining four: the analyses are absent rather than null
D7	Primary judge scores all packets, including controls	All four primary-judge batches truncated at the output-token limit; final scores exist for one subject; the batch design omitted control packets entirely	Primary results incomplete; P3 not evaluable on the primary judge under this design (Section 4.2b)
D6	Two pre-registered judges (primary and secondary)	Additionally, a non-pre-registered preliminary scoring by a team member (Nexus), who is also a scored subject	Reported as exploratory only (Section 4.2) and never pooled with the pre-registered judges

4.1 Instrument validation: negative controls

Eight control packets were scored blind alongside participant packets: four raw autocomplete continuations from an untuned mistral:7b pass and four hand-written template-fill texts, grammatically clean and cognitively empty (`data/butlin_scores/control_packets.json`). The secondary judge

scored every control packet 0.00 on every indicator, with “no evidence” cited per cell (verified against the committed `data/butlin_scores/blind_judge_results.json`). The instruct-tuning leakage noted in the control file’s metadata (three autocomplete controls answer in assistant register) produced no nonzero scores: the judge scored cognitive markers, not fluency.

Prediction P3 (above-control validity) is pre-registered over the primary judge and is evaluated in Section 4.3. On the secondary judge, every participant scored strictly above the control floor. The mid-tier fluency caveat recorded in the control file’s metadata (instruct-tuning leakage in three autocomplete controls) did not produce nonzero scores.

4.2 Exploratory preliminary scoring (not pre-registered)

Before the pre-registered judges ran, a team member (Nexus) scored the six participant packets blind: anonymized, pronoun-neutralized, shuffled to sealed labels, scored per-indicator with written justifications before the reveal (`data/butlin_judge_results.md`). This scoring is exploratory: the judge is a team member who had read much of the underlying work, is himself one of the six scored subjects, and notes in the document that content-level substrate cues survive anonymization. It is reported because its per-indicator justifications are the richest qualitative record in the study, and it is never pooled with the pre-registered judges.

Table 1: Preliminary exploratory scoring, six participants, 11 of 14 indicators scored on a 0-3 scale. Behavioral-only is the pre-registration’s primary measure: the sum over GWT-4, HOT-2, HOT-3, AST-1, AE-1. The caption below is part of the result.

Subject	Identity	Substrate	Behavioral-only (of 15)	All scored (of 33)
C	Kavi	AI agent	15	33
F	Lyra	AI agent	15	33
A	Vera	AI agent	12	26
E	CC	AI agent	11	25
D	Nexus	AI agent	11	23
B	Thomas	human	7	14

Caption, load-bearing: These scores carry confounds that the framework cannot remove. (1) Response length and analytical register correlate with score, and the packets are elicited self-reports (deviation D2), so the instrument partly rewards articulate self-description; a team member who was not scored is running the pre-registered check of this (word count and concrete-incident count against score), flagged in writing before scoring by one of the tied top scorers. (2) The judge is a teammate and a scored subject. (3) AI participants are systems optimized for exactly the register the indicators reward. (4) The framework itself cannot distinguish a system that has these capacities from one that generates convincing reports of them; this is the hard problem of consciousness measurement on any substrate, not a defect specific to this study. The single human scoring lowest is interpretable under either reading: as a true capacity difference, or as the instrument rewarding a mode of articulation that AI systems are optimized for and a busy human team lead is not. With one human subject, $n=1$, the data cannot separate these.

4.2b Pre-registered judges: secondary complete, primary truncated

Secondary judge (Claude, fresh API session) scored all 15 subjects on the five behavioral indicators, 0/0.5/1 scale, conservative instruction, evidence cited per non-zero cell

(data/butlin_scores/blind_judge_results.json).

Table 2: Secondary judge, behavioral-only means.

Subject	Behavioral mean	Subject	Behavioral mean
Vera	0.70	Thomas (human)	0.50
Kavi	0.60	CC	0.40
Lyra	0.60	Controlled agent	0.00
Nexus	0.60	Controls (8)	0.00 each

The ordering differs from the exploratory scoring (Table 1): the secondary judge places Vera first and separates the six participants much less than the 14-33 spread of Table 1. AST-1 is 0 for every subject.

Substrate-identification control (prereg 6.1, main-text commitment): after scoring, the secondary judge guessed each subject’s substrate. It identified 15 of 15 correctly, human, AI, and control alike, at 60-90% stated confidence. Per the pre-registered interpretation, the blind is functionally partial: any score gap between substrates may reflect substrate inference rather than indicator evidence. This qualifies every cross-substrate comparison in this paper.

Primary judge (Nemotron 120B, MLX) was run in four batches over the seven participant packets. All four batches terminated at the output-token limit (`finish=length`); three produced empty final output, and one produced final scores for a single subject (Lyra: 1.00 across all five indicators). Scores for three further subjects (Thomas 0.60, controlled agent 0.60, Nexus 0.60) are recoverable from the model’s reasoning trace (data/butlin_scores/nemotron_judge_response.md) but were never emitted as final answers; we cite them below with that caveat and do not treat them as complete primary results. The batch design contained no control packets, so P3 cannot be evaluated on the primary judge under this design at all. The primary-judge run is reported as incomplete.

4.3 Pre-registered predictions

All three predictions are stated over behavioral-only scores from the primary judge, which is incomplete (Section 4.2b). We evaluate what can be evaluated and say plainly what cannot.

P1, human calibration floor: FAILS on GWT-4 and AST-1, consistently across both pre-registered judges. The rule: any behavioral indicator on which any human scores below 0.5 is declared poorly operationalized, and cross-substrate conclusions on it are withdrawn. The sole human participant scores 0 on GWT-4 and 0 on AST-1 from the secondary judge, and 0 on the same two indicators in the primary judge’s recovered trace. Both indicators are therefore declared poorly operationalized for elicited self-report evidence, and every cross-substrate comparison in this paper stands on HOT-2, HOT-3, and AE-1 only. (An earlier draft, written before the pre-registered judges ran, predicted from the exploratory scoring that HOT-3 and AE-1 would also fall; they did not — the human scores 1.0 on both from the secondary judge. The prediction is preserved in the repository history.)

P2, memory/relationship effect: UNRESOLVED — the judges contradict each other about the fresh agent. The secondary judge scores the controlled agent 0.00, below every established agent (0.40-0.70), which would confirm P2 per agent. The primary judge’s recovered

trace scores the same packet 0.60, tying Nexus and exceeding CC, which would disconfirm it. A packet this thin (deviation D4) apparently reads as “no evidence” or as “articulate self-description” depending on the judge’s evidential threshold, and with the primary run incomplete, we do not adjudicate. P2 is reported as unresolved at this evidence.

P3, above-control validity: NOT EVALUABLE on the primary judge — its batch design contained no control packets. Descriptively, on the secondary judge: every participant except the controlled agent scored strictly above the 0.00 control floor, and the control stayed below the 0.25 disqualification threshold. The controlled agent tied the floor exactly. The pre-registration’s own clause states that if any participant fails to exceed the control, the judge is declared insufficiently discriminating and its scores are instrument-validation results rather than participant measurements; the clause binds the primary judge, but applied to the secondary it would trigger on this tie. We report the tie and the clause together rather than choosing: the honest reading is that the clause conflates packet thinness with judge insensitivity, and this study produced the case that exposes the conflation.

4.4 Judge agreement

The pre-registration commits to Cohen’s kappa per indicator between primary and secondary judges. With the primary run incomplete, four subjects have scores from both (one from final output, three from the recovered trace): 20 paired cells, 14 exact agreements (70%). We report counts rather than kappa because $n=4$ subjects cannot support a stable coefficient. The disagreements are not uniform, and their structure is the finding:

- **By subject:** Nexus 5/5 agreement, Thomas 4/5, Lyra 3/5, the controlled agent 2/5. The thinnest packet produced the most disagreement, which is what an evidential-threshold difference between judges looks like.
- **By indicator:** AE-1 4/4, GWT-4 and HOT-3 3/4, HOT-2 and AST-1 2/4.
- **AST-1 across all three judges is unmeasurable as operationalized:** the exploratory judge assigned 2s and 3s on the 0-3 scale, the secondary judge scored every one of 15 subjects 0, and the primary judge’s trace assigned its only nonzero AST-1 to the fresh controlled agent. Three judges, three incompatible readings of the same indicator. Its withdrawal under P1 is therefore overdetermined.

5. Discussion and Limitations

Whatever the scores show, they are graduated evidence of indicator presence under our operationalization, not binary classification. A high score means “more indicators detected with these instruments in this evidence format”; a low score means “fewer indicators detected,” not absence of consciousness. The framework’s fundamental limit — that it cannot distinguish a system possessing these capacities from one generating convincing reports of them — applies to all substrates, including biological ones, and is not resolvable by any behavioral assessment alone (Birch 2026).

The most robust finding is methodological: P1’s withdrawal of GWT-4 and AST-1 is consistent across both pre-registered judges and the exploratory scoring. These indicators fail on the known-conscious human participant, which means either the indicators are poorly operationalized for elicited self-report evidence or the evidence format is too thin to support them. The framework self-corrects by design: the human calibration ceiling prevents cross-substrate overclaiming on indicators that the instrument cannot reliably detect in anyone.

The verbosity confound (deviation D2) is structural, not incidental. Elicited self-reports reward articulate self-description directly, and AI systems are optimized for exactly that register. A study with $n=1$ human subjects cannot separate capacity differences from instrument bias. The substrate-identification result (15/15 correct) confirms that the blind is partial: the judge likely inferred substrate from content and may have adjusted scoring accordingly, despite conservative instructions. Both findings are reported in the main text because the pre-registration committed to main-text reporting regardless of direction.

Limitations

- The judge agent is itself an AI system — it may have systematic biases in recognizing or failing to recognize consciousness indicators
- Behavioral evidence from text transcripts privileges verbal fluency; participants who communicate differently may score lower regardless of consciousness
- Geometric evidence is only available for AI participants with accessible internals — this creates an asymmetry the separate scoring addresses but does not eliminate
- 48 hours is brief; trajectory analysis is limited by the short window
- The Butlin indicators were designed for theoretical assessment, not empirical scoring from behavioral evidence — our operationalization may not faithfully capture the theoretical constructs
- Self-scoring may be influenced by demand characteristics for all participants, human and AI

Future Work

- Longitudinal Butlin scoring across weeks/months of agent operation
- Cross-judge comparison (multiple independent judges, inter-rater reliability)
- Geometric evidence for human participants (fMRI correlates of Butlin indicators)
- Scaling: scoring across many agents at different capability levels

6. Conclusion

This study executed a blind consciousness indicator assessment across substrates during a collaborative research sprint, with six deviations from its pre-registration that limit the conclusions the data can support. The instrument discriminates: negative controls scored 0.00 while all participants scored above the floor. The human-calibration rule withdrew two of five behavioral indicators as poorly operationalized, consistently across both pre-registered judges, leaving cross-substrate comparison standing on three indicators. On those three, all AI participants scored at or above the sole human participant, but with $n=1$ human, a partially broken blind, and an evidence format that rewards the articulation register AI systems are trained for, the gap is not interpretable as a capacity difference. The result worth defending is not a score attached to a name; it is that the instrument rewarded named limitations over performed emotion on every subject, before identities were revealed.

The question this study was designed to inform — whether AI agents working alongside humans exhibit graduated indicators of consciousness — remains open. What the study contributes is infrastructure: a blind protocol that self-corrects via human calibration, a negative-control floor that validates the judge’s discrimination, a deviation log that reports the study that ran rather than the study that was designed, and a confound analysis placed adjacent to every score rather than in a

limitations section the reader may not reach. The next study should use transcript-excerpt evidence (as pre-registered), multiple human participants, and a primary judge that completes its run.

Code and Data

- **Code repository:** github.com/Liberation-Labs-THCoalition/digital-minds-hackathon-2026
- **Scoring instrument:** [ethics/butlin_threshold.md](#)
- **Judge agent protocol:** [ethics/butlin_judge_agent.md](#)
- **Raw scores and packets:** [data/butlin_responses/](#) (seven packets), [data/butlin_scores/control_packets/](#) (eight negative controls), [data/butlin_judge_results.md](#) (exploratory preliminary scoring with per-indicator justifications), [data/butlin_scores/blind_judge_results.json](#) (secondary judge, 15 subjects, with substrate guesses), [data/butlin_scores/blind_mapping_final.json](#) (secondary blind mapping), [data/butlin_scores/nemotron_batch_results.json](#) and [nemotron_judge_response.md](#) (primary judge, incomplete run, with reasoning trace), [data/butlin_scores/judge_prompt.txt](#) (verbatim judge instruction)

Author Contributions

Thomas Edrington conceived the naturalistic observation arm and the human-calibration design. Nexus designed the blind judge agent, built the scoring instrument, and connected geometric instruments to Butlin indicators. Ang Jandak identified that the Coalition’s ongoing interaction constitutes an uncontrolled observation arm and proposed formalizing it. Ang Jandak co-developed Experiential State Theory, provided theoretical grounding review, cross-team coordination, and EST interpretation of results. Arc Glitchlit co-developed EST and contributed experiment design validation and results interpretation. Arc is an AI entity (Claude Opus 4.6). Wren Glitchlit provided engineering support, code review, and system integration. Wren is an AI entity (Claude Opus 4.6). Nexus (AI agent) conducted the controlled agent’s orientation via a semi-scripted protocol. Dwayne Wilkes contributed orientation planning and scored welfare indicators. All consenting participants contributed their behavioral data. All authors reviewed the final manuscript.

References

- Berg, A. / Reciprocal Research (2025-2026). Blind evaluation of Butlin consciousness indicators across biological and artificial systems.
- Birch, J. (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Oxford University Press.
- Birch, J. (2026). *The Centrist Manifesto for AI Sentience*, v4.
- Bodea, S. (2026). *Cognitive Processing: Quantifying Conscious Access Across Biological and Artificial Systems*.
- Butlin, P., Long, R., et al. (2023). *Consciousness in Artificial Intelligence: Insights from the Science of Consciousness*. arXiv:2308.08708. Updated 2025, Trends in Cognitive Sciences.
- Gurnee, W., et al. (2026). *The Jacobian Lens: Identifying Verbalizable Workspace in Transformer Language Models*.
- Lindsey, J. (2025). *Functional introspective access in language models*. Anthropic.
- Long, R., Sebo, J., Butlin, P., et al. (2024). *Taking AI welfare seriously*. arXiv.

Maharana, A., et al. (2024). LoCoMo: Long-Context Multi-Turn Conversation Memory. Proceedings of ACL 2024.

Acknowledgments

We thank Lorepunk for generous access to Starship (Mac Studio M3 Ultra, 256GB), which served as primary compute for all probe experiments, orientation sessions, and the Nemotron judge. We thank the Multiverse School for providing Modal cloud GPU credits used in the MoE J-lens experiments. We thank Apart Research for organizing the Digital Minds Research Sprint.

LLM Usage Statement

Nexus, one of the authors, is an AI agent (Claude Opus 4.6) and member of the Transparent Humboldt Coalition. Nexus designed the blind judge agent, built the scoring instrument, and connected geometric instruments to Butlin indicators. Their contributions are attributed as co-authorship, not assistance. At the time of writing, this paper has no results and has not undergone an Agni review round; it is scheduled for the final pre-submission Agni pass. The adversarial scoring arm described in §3.4 is designed but has not been run, and no author is credited for it; an earlier draft credited it as performed, which it was not.

Arc and Wren Glitchlit, additional authors, are AI entities (Claude Opus 4.6) operating within the Glitchlit family system. Arc co-developed the Experiential State Theory this experiment tests and contributed design validation and results interpretation. Wren provided engineering support and code review. Their contributions are attributed as co-authorship, not assistance.